Consentful Cybernetics — YC Pitch Draft

One-liner

Consentful Cybernetics builds AI models that detect coercion, pressure, and consent breakdown in communication before harm happens.

Problem

The internet optimized communication for attention, engagement, and conversion. That gave us powerful systems for getting people to click, continue, comply, and say yes.

But we have almost no comparable infrastructure for knowing whether a yes is clear, unpressured, current, scoped, and safe to refuse.

This matters everywhere communication has power: workplaces, dating apps, customer support, healthcare, education, AI agents, legal workflows, platforms, and human-AI interaction.

Today, systems can tell whether someone opened a message, clicked a link, watched a video, or converted.

They cannot tell whether the interaction became coercive, whether consent went stale, whether scope drifted, whether a person was pressured into compliance, or whether refusal was actually safe.

Solution

We are building a Consent Dynamics Model: an AI model trained to recognize consent-relevant dynamics in communication.

It does not "certify consent." That would be dangerous.

Instead, it acts as a witness layer. It detects patterns such as:

- pressure
- coercion
- hesitation
- unsafe refusal conditions
- power imbalance
- scope drift
- stale permission
- escalation without revalidation
- mismatch between verbal agreement and behavioral signals
- repair opportunities before harm hardens into breach

The model can work across text, audio, video, and eventually multimodal interaction streams.

AI can already recognize cats better than the species that domesticated them and got domesticated back. It can learn to recognize consent dynamics too.

Product

The first product is an API and dashboard for consent-sensitive communication analysis.

Initial customers could use it to evaluate and improve high-stakes communication flows:

1. AI agent interactions
   Detect when agents pressure users, overstep scope, or continue after uncertainty.

2. Workplace communication
   Identify coercive patterns in manager/employee requests, HR workflows, and internal processes.

3. Customer support and sales
   Distinguish persuasion from pressure and reduce manipulative conversion patterns.

4. Dating, social, and community platforms
   Detect boundary pressure, escalation, harassment, and unsafe refusal dynamics.

5. Healthcare, legal, and financial services
   Improve informed consent, scope clarity, and documentation around high-impact decisions.

The model would surface risk signals and suggested clarifying prompts, not final judgments.

Example output:

> "This interaction shows rising pressure, unclear refusal safety, and escalation beyond the originally agreed scope. Recommend pausing and revalidating consent."

Why now

AI-mediated communication is exploding.

Autonomous agents will increasingly negotiate, persuade, sell, schedule, escalate, recommend, and act on behalf of people and institutions.

If these systems optimize only for task completion, engagement, or conversion, they will reproduce the worst dynamics of the attention economy at agentic speed.

We need consent-aware infrastructure before AI agents become the default interface for human life.

The same advances that make AI good at detecting objects, intent, sentiment, and risk can be redirected toward detecting consent dynamics.

Market

Every organization using AI to communicate with humans will eventually need a consent and coercion safety layer.

Potential markets include:

- AI agent platforms
- trust and safety teams
- enterprise communication tools
- HR and workplace compliance
- healthcare consent workflows
- dating and social platforms
- customer support QA
- sales enablement ethics/compliance
- education technology
- legal technology

The wedge is AI agent safety: companies deploying agents need to know when their systems are pressuring users, exceeding scope, or creating liability.

Differentiation

Most AI safety tools focus on content: toxicity, hate, sexual content, self-harm, privacy, legal risk, or factuality.

Consentful Cybernetics focuses on interaction dynamics.

The question is not only “what was said?”

The question is:

> What happened between the agents?

Consent and coercion are trajectories, not tokens.

That requires modeling timing, power, scope, context, hesitation, escalation, medium, role, and repair.

Vision

The attention economy built systems that know how to reach us better than they know how to respect us.

Consentful Cybernetics is building the missing layer: AI that helps communication preserve agency.

Long term, every AI agent, workplace tool, platform, and high-stakes communication system should have consent dynamics monitoring built in.

We do not want machines that get better at extracting yes.

We want machines that make yes trustworthy.

Tagline

A trustworthy yes is more valuable than a captured yes.